

Reinforcement Learning for Shanghai Stock Index Trading Strategy

Abstract

This report presents a reinforcement learning approach to develop automated trading strategies for the Shanghai Stock Exchange Composite Index. We implement and compare two state-of-the-art policy gradient algorithms—**A2C (Advantage Actor-Critic)** and **PPO (Proximal Policy Optimization)**—within a custom trading environment that complies with China's T+1 trading regulation. Our experiments demonstrate that the A2C algorithm achieves a **10.92% return** on the 2024-2025 test period, outperforming both the PPO strategy (8.77%) and a random baseline (2.21%).

1. Introduction and Background

1.1 Problem Context

The stock market presents a challenging sequential decision-making problem where traders must decide when to buy, hold, or sell assets based on historical price patterns and technical indicators. Traditional approaches rely on rule-based strategies or fundamental analysis, which often fail to capture complex nonlinear market dynamics.

The **Shanghai Stock Exchange Composite Index** (SSE Composite) is one of the most important financial indicators in Asia, tracking all stocks listed on the Shanghai Stock Exchange. Trading in Chinese A-shares is governed by the **T+1 settlement rule**, which prohibits selling shares on the same day they were purchased. This is a critical constraint that must be incorporated into any realistic trading model.

1.2 Project Objectives

The primary goals of this project are:

1. **Design a compliant trading environment** that accurately models the T+1 trading constraints of the Chinese stock market
2. **Implement technical indicator features** to provide rich market state representations
3. **Train and compare RL algorithms** (A2C and PPO) for automated trading decisions
4. **Evaluate performance** on out-of-sample data (2024-2025) to assess generalization capability

1.3 Related Work

Reinforcement learning has gained significant traction in financial trading applications. Notable frameworks include:

- **gym-anytrading**: An OpenAI Gym environment for stock trading simulation
- **Stable-Baselines3**: A reliable implementation of RL algorithms in PyTorch
- **FinRL**: A library specifically designed for financial reinforcement learning

Our work extends these foundations by implementing T+1 compliance and eliminating look-ahead bias in feature computation.

2. Mathematical Models and Formulas

2.1 Reinforcement Learning Framework

We formulate the trading problem as a **Markov Decision Process (MDP)** defined by the tuple (S, A, P, R, γ) :

Component	Symbol	Description
State Space	S	Market observations including technical indicators
Action Space	$A = \{0,1\}$	Buy (0) or Sell (1)
Transition Probability	$P(s' s, a)$	Market dynamics (stochastic)
Reward Function	$R(s, a, s')$	Trading profit/loss
Discount Factor	$\gamma = 0.99$	Future reward weighting

2.2 Technical Indicators as State Features

The state representation consists of 10 normalized technical indicators computed over a 30-day sliding window.

2.2.1 Moving Averages (MA)

The n-period moving average at time t is calculated as:

$$MA_n(t) = \frac{1}{n} \sum_{i=0}^{n-1} P_{t-i}$$

where P_t is the closing price at time t . We compute MA5, MA10, MA20, and MA60.

The MA Deviation Ratio measures the relative distance between current price and moving average:

$$MA_n\text{-Ratio}(t) = \frac{P_t - MA_n(t)}{MA_n(t)} \times 100$$

2.2.2 Relative Strength Index (RSI)

The RSI indicator is computed as:

$$RSI = 100 - \frac{100}{1 + RS}$$

where RS (Relative Strength) is:

$$RS = \frac{\text{Average Gain over } n \text{ periods}}{\text{Average Loss over } n \text{ periods}}$$

Using Wilder's smoothing method with $n = 14$:

$$\text{Avg Gain}_t = \frac{\text{Avg Gain}_{t-1} \times (n - 1) + \text{Gain}_t}{n}$$

2.2.3 Moving Average Convergence Divergence (MACD)

The MACD indicator consists of three components:

$$\text{MACD Line} = \text{EMA}_{12} - \text{EMA}_{26}$$

$$\text{Signal Line} = \text{EMA}_9(\text{MACD Line})$$

$$\text{Histogram} = \text{MACD Line} - \text{Signal Line}$$

The Exponential Moving Average (EMA) is defined as:

$$\text{EMA}_t = P_t \cdot \alpha + \text{EMA}_{t-1} \cdot (1 - \alpha), \quad \alpha = \frac{2}{n + 1}$$

2.2.4 Bollinger Bands

Bollinger Bands provide volatility-based price channels:

$$\text{Upper Band} = \text{MA}_{20} + 2\sigma$$

$$\text{Lower Band} = \text{MA}_{20} - 2\sigma$$

$$\text{BB Position} = \frac{P_t - \text{Lower Band}}{\text{Upper Band} - \text{Lower Band}}$$

where σ represents the 20-period standard deviation of closing prices.

2.3 Policy Gradient Algorithms

2.3.1 Advantage Actor-Critic (A2C)

A2C learns both a policy $\pi_\theta(a|s)$ (actor) and a value function $V_\phi(s)$ (critic). The advantage function reduces variance in policy gradient estimation:

$$A(s_t, a_t) = R_t + \gamma V_\phi(s_{t+1}) - V_\phi(s_t)$$

The policy is updated using the gradient:

$$\nabla_\theta J(\theta) = \mathbb{E}[\nabla_\theta \log \pi_\theta(a_t|s_t) \cdot A(s_t, a_t)]$$

The value function is trained to minimize the temporal difference error:

$$L_V(\phi) = \mathbb{E}[(R_t + \gamma V_\phi(s_{t+1}) - V_\phi(s_t))^2]$$

2.3.2 Proximal Policy Optimization (PPO)

PPO improves training stability through a clipped surrogate objective:

$$L^{CLIP}(\theta) = \mathbb{E}[\min(r_t(\theta)A_t, \text{clip}(r_t(\theta), 1 - \epsilon, 1 + \epsilon)A_t)]$$

where the probability ratio is:

$$r_t(\theta) = \frac{\pi_\theta(a_t|s_t)}{\pi_{\theta_{old}}(a_t|s_t)}$$

The clip parameter $\epsilon = 0.2$ prevents excessively large policy updates, ensuring stable learning.

2.4 Reward Function Design

The reward function considers both realized profits and unrealized positions:

Trade Reward (on sell):

$$R_{trade} = \frac{P_{sell} - P_{buy}}{P_{buy}} \times 100$$

Holding Reward:

$$R_{hold} = 0.1 \times \frac{P_t - P_{t-1}}{P_{t-1}} \times 100$$

Transaction Costs:

Cost Type	Rate
Buy commission	0.03%
Sell stamp duty	0.10%

3. Implementation Details

3.1 Data Description

Attribute	Value
Data Source	Shanghai Stock Exchange Composite Index
Date Range	1990-12-19 to 2026-01-09
Total Samples	8,560 trading days
Features	OHLCV + Technical Indicators

Data Split:

Dataset	Period	Samples
Training	1991-01-31 to 2023-12-29	8,040
Testing	2024-01-02 to 2026-01-09	490

3.2 Environment Architecture

The custom trading environment (ShanghaiIndexEnvT1) is built upon the gym-anytrading framework with the following specifications:

Component	Specification
Observation Space	Box(30, 10) — 30-day window × 10 features
Action Space	Discrete(2) — Buy/Hold (0) or Sell (1)
T+1 Compliance	Tracks holding days; blocks same-day selling
Look-ahead Prevention	Decisions use only T-1 and earlier data

The environment enforces the T+1 trading rule by maintaining a holding day counter. After a buy action, the agent must hold the position for at least one trading day before selling is permitted.

3.3 Feature Engineering

Index	Feature	Normalization	Range
1	Daily Price Change	Percentage	[-100, 100]
2	MA5 Deviation	Percentage	[-100, 100]
3	MA10 Deviation	Percentage	[-100, 100]
4	MA20 Deviation	Percentage	[-100, 100]
5	MA5-MA10 Cross	Percentage	[-100, 100]
6	RSI (14-day)	Min-max scaled	[0, 1]
7	MACD Histogram	Max-normalized	[-1, 1]
8	Bollinger Position	Min-max scaled	[0, 1]
9	Daily Amplitude	Percentage	[0, 100]
10	Candlestick Body	Percentage	[-100, 100]

3.4 Neural Network Architecture

Both A2C and PPO employ identical Multi-Layer Perceptron (MLP) architectures:

Layer	Configuration
Input Layer	300 neurons (30 × 10 flattened)
Hidden Layer 1	256 neurons with ReLU activation
Hidden Layer 2	256 neurons with ReLU activation

Layer	Configuration
Policy Head	2 outputs (action probabilities)
Value Head	1 output (state value estimate)

3.5 Training Configuration

Hyperparameter	A2C	PPO
Learning Rate	0.0007	0.0003
Steps per Update	5	2,048
Batch Size	—	64
Epochs per Update	—	10
Discount Factor γ	0.99	0.99
Entropy Coefficient	0.01	0.01
Value Function Coefficient	0.5	—
Clip Range ϵ	—	0.2
Total Timesteps	100,000	100,000

3.6 Eliminating Look-ahead Bias

A critical implementation detail is preventing look-ahead bias in the observation function. When making decisions at time T , the agent only has access to data from $T-1$ and earlier periods. The current day's data (time T) is used solely for trade execution, not for decision-making. This design ensures that the backtesting results reflect realistic trading conditions.

4. Results and Performance Evaluation

4.1 Performance Summary

Strategy	Avg Return	Return Rate	Max	Min	Std Dev
Random	1.0221	+2.21%	1.2579	0.7904	0.1265
A2C	1.1092	+10.92%	1.1092	1.1092	0.0000
PPO	1.0877	+8.77%	1.0877	1.0877	0.0000

Note: A2C and PPO show zero variance because deterministic evaluation produces identical results across episodes.

4.2 Trading Statistics

Metric	Random	A2C	PPO
Number of Trades	127	98	120
Buy Signals	127	98	120
Sell Signals	126	98	120
Final Profit	+4.01%	+10.92%	+8.77%

4.3 Visual Analysis

Figure 1 illustrates the return distribution comparison across the three strategies. The random strategy exhibits high variance with returns ranging from -20.96% to +25.79%, while both A2C and PPO produce deterministic, consistent returns. The A2C strategy achieves the highest return at 10.92%, followed by PPO at 8.77%.




Figure 1: Performance Comparison - Return distribution (left) and average returns (right) across strategies

Figure 1: Performance comparison showing return distribution (left) and average returns with standard deviation (right) for Random, A2C, and PPO strategies on the 2024-2025 test period.

4.4 Trading Visualization

The following figures show the trading signals generated by each strategy on the test period. Green triangles indicate buy signals, red triangles indicate sell signals, and dashed lines connect each buy-sell pair with color indicating profit (green) or loss (red).




Figure 2: Random Strategy Trading Signals

Figure 2: Random strategy trading visualization showing buy/sell signals on the Shanghai Index (2024-2025).




Figure 3: A2C Strategy Trading Signals

Figure 3: A2C strategy trading visualization showing learned buy/sell signals with 10.92% total return.




Figure 4: PPO Strategy Trading Signals

Figure 4: PPO strategy trading visualization showing learned buy/sell signals with 8.77% total return.

4.5 Sample Trade Analysis (A2C)

Selected profitable trades from the A2C strategy demonstrate its ability to capture market movements:

Trade #	Buy Price	Sell Price	Return	Hold Days
2	2,831	3,005	+6.16%	12
23	2,861	3,001	+4.90%	36

Trade #	Buy Price	Sell Price	Return	Hold Days
15	3,045	3,148	+3.38%	6
25	3,169	3,262	+2.91%	1

4.6 Algorithm Comparison

Aspect	A2C	PPO
Training Time	~2 min	~2.5 min
Sample Efficiency	Higher	Lower
Policy Stability	Moderate	High
Final Return	10.92%	8.77%
Trade Frequency	Moderate (98)	Higher (120)

Key Finding: A2C's more frequent and smaller policy updates prove more effective for this financial application, achieving better returns with fewer but more strategic trades compared to PPO.

5. Discussion

5.1 Key Findings

1. **A2C outperforms PPO** by a margin of 2.15 percentage points, demonstrating that the simpler advantage actor-critic approach can be more effective for this financial application.
2. **Both RL methods significantly outperform random trading**, validating that the learned policies capture meaningful market patterns.
3. **T+1 compliance integration** successfully prevents unrealistic intraday round-trips that would inflate backtesting returns.
4. **Technical indicators provide discriminative features** for the RL agents to learn trend-following and mean-reversion patterns.

5.2 Market Implications

The trained models demonstrate capabilities to:

- **Identify trend reversals** using MACD and MA crossover signals
- **Avoid overtrading** during sideways markets
- **Capture momentum** during strong directional moves

5.3 Limitations

Limitation	Description	Impact
Position Sizing	Binary buy/sell actions only	Cannot optimize

Limitation	Description	Impact
Single Asset	Index only, no diversification	portfolio allocation Concentrated risk exposure
Transaction Costs	Simplified fee structure	May underestimate real trading friction
Market Impact	Assumes infinite liquidity	Not applicable for large capital
Non-stationarity	Market regime changes	Potential performance decay over time
Look-ahead in Indicators	Some technical indicators use future data implicitly	May cause minor data leakage

5.4 Future Work

1. **Multi-asset Portfolio Optimization:** Extend to multiple stocks with correlation modeling
2. **Continuous Action Space:** Implement portfolio weight allocation
3. **Advanced Architectures:** Explore LSTM/Transformer for sequential modeling
4. **Risk-adjusted Rewards:** Incorporate Sharpe ratio or maximum drawdown constraints
5. **Online Learning:** Implement incremental learning for regime adaptation

6. Conclusion

This project successfully demonstrates the application of deep reinforcement learning to Shanghai Stock Index trading under realistic T+1 constraints. The **A2C algorithm achieves a 10.92% return** on the 2024-2025 test period, outperforming both the PPO strategy (8.77%) and random trading baseline (2.21%).

Key contributions include:

- A **T+1 compliant trading environment** that accurately models Chinese market regulations
- **Elimination of look-ahead bias** through careful feature engineering
- **Comprehensive evaluation** comparing multiple RL algorithms

While the results are promising, real-world deployment would require additional considerations including position sizing, risk management, and robustness testing across different market regimes.

References

- [1] Schulman, J., Wolski, F., Dhariwal, P., Radford, A., & Klimov, O. (2017). Proximal Policy Optimization Algorithms. *arXiv preprint arXiv:1707.06347*.
- [2] Mnih, V., Badia, A. P., Mirza, M., Graves, A., Lillicrap, T., Harley, T., ... & Kavukcuoglu, K. (2016). Asynchronous Methods for Deep Reinforcement Learning. *International Conference on Machine Learning (ICML)*.
- [3] OpenAI Gymnasium Documentation. Available at: <https://gymnasium.farama.org/>
- [4] Stable-Baselines3 Documentation. Available at: <https://stable-baselines3.readthedocs.io/>
- [5] gym-anytrading: Stock Trading Environment for OpenAI Gym. Available at: <https://github.com/AminHP/gym-anytrading>
-

Appendix

A. Environment Parameters

Parameter	Value
Window Size	30 days
Random Seed	42
Total Training Timesteps	100,000
Buy Commission	0.03%
Sell Stamp Duty	0.10%
Hidden Layer Dimensions	[256, 256]
Activation Function	ReLU

B. Feature List Summary

Feature	Description
pct_change	Daily return percentage
ma5_ratio	Price deviation from 5-day MA
ma10_ratio	Price deviation from 10-day MA
ma20_ratio	Price deviation from 20-day MA
ma5_10_cross	MA5-MA10 crossover signal
rsi_norm	Normalized RSI (14-day)
macd_norm	Normalized MACD histogram
bb_position	Position within Bollinger Bands
amplitude	Daily price amplitude
body	Candlestick body size

C. Reproducibility

To ensure reproducibility, the following random seeds were set:

Library	Seed Value
PyTorch	42
NumPy	42
Python random	42

Required packages: Python 3.x, PyTorch, Stable-Baselines3, gym-anytrading, pandas, numpy, matplotlib.

Report: January 2026

Test Period: 2024-01-02 to 2026-01-09

Written By: Yongheng Wang, Ping-Jui Chen